

ER-Relational Mapping (cont)

1/28

So far, have considered mappings for ...

- ER attribute \rightarrow relational attribute
- ER entity \rightarrow relational table
- ER key \rightarrow primary key for table
- $n:m$ relationship \rightarrow relational table
(with foreign key for each participating entity plus relationship attributes)
- $1:n$ relationship \rightarrow foreign key plus relationship attributes
- $1:1$ relationship \rightarrow foreign key plus relationship attributes

n-way Relationships

2/28

Relationship mappings above assume binary relationship.

If multiple entities are involved:

- $n:m$ generalises naturally to $n:m:p:q$
 - include foreign key for each participating entity
 - include any other attributes of the relationship
- other multiplicities (e.g. $1:n:m$) ...
 - need to be mapped the same as $n:m:p:q$
 - so not quite an accurate mapping of the ER

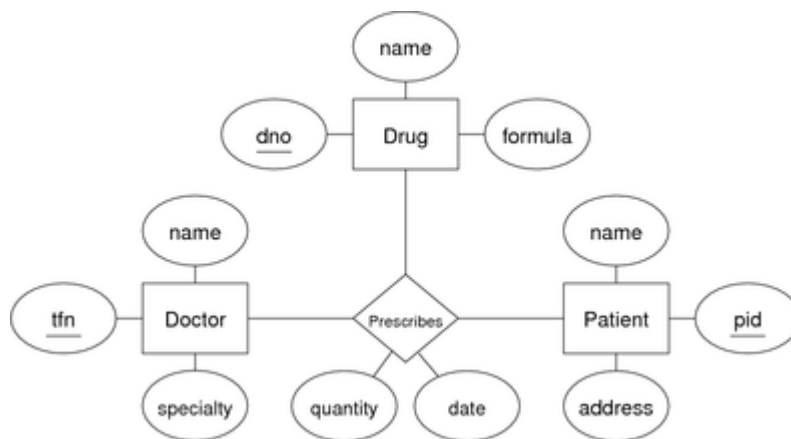
Some people advocate converting n-way relationships into:

- a new entity, and a set of n binary relationships

Exercise: 3-way relationship

3/28

Translate the following ER design to a relational schema:

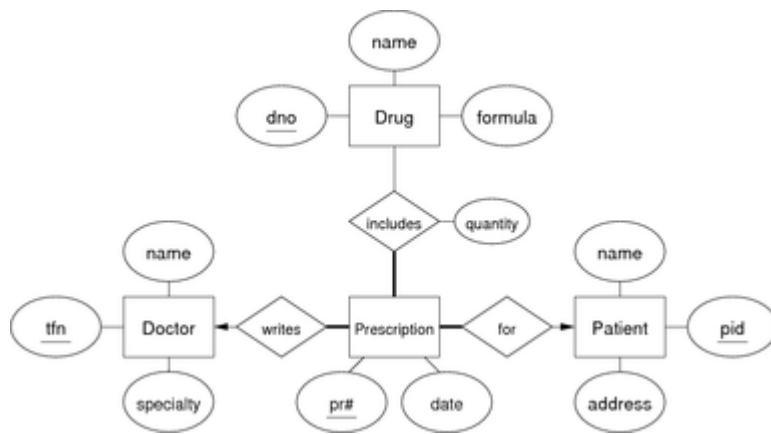


[Solution]

Exercise: Alternative prescription model

4/28

Translate the following ER design to a relational schema:



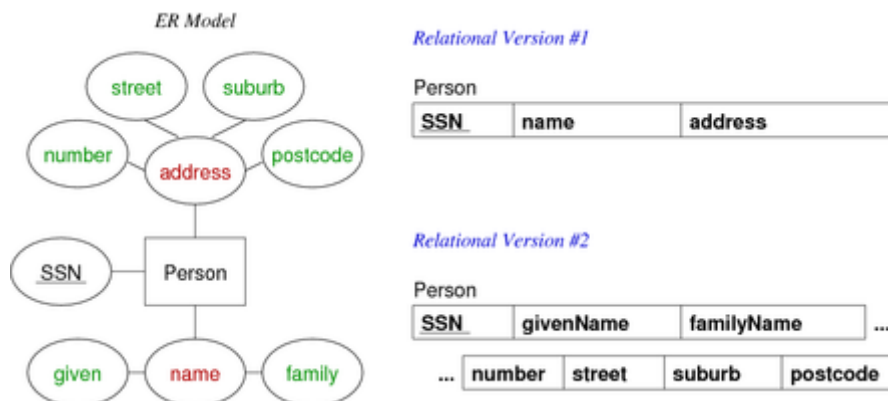
[Solution]

Mapping Composite Attributes

5/28

Composite attributes are mapped by concatenation or flattening.

Example:

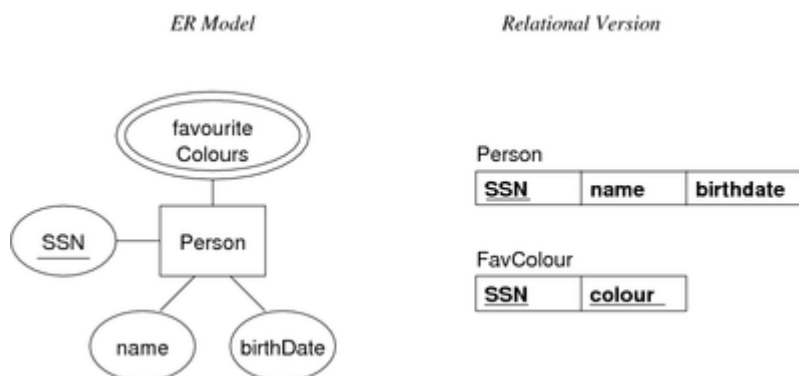


Mapping Multi-valued Attributes (MVAs)

6/28

MVAs are mapped by a new table linking values to their entity.

Example:



... Mapping Multi-valued Attributes (MVAs)

7/28

Example: the two entities

```
Person(12345, John, 12-feb-1990, [red,green,blue])
Person(54321, Jane, 25-dec-1990, [green,purple])
```

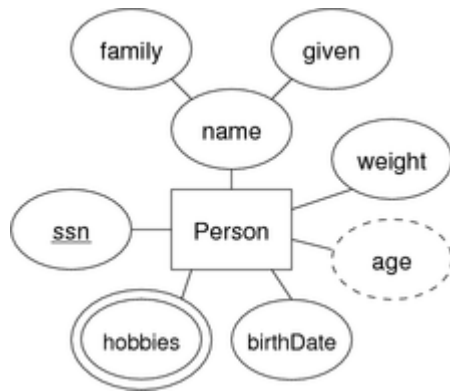
would be represented as

Person(12345, John, 12-feb-1990)
Person(54321, Jane, 25-dec-1990)
FavColour(12345, red)
FavColour(12345, green)
FavColour(12345, blue)
FavColour(54321, green)
FavColour(54321, purple)

Exercise: Attribute Mappings

8/28

Convert this ER design to relational form:



[\[Solution\]](#)

Mapping Subclasses

9/28

Three different approaches to mapping subclasses to tables:

- ER style
 - each entity becomes a separate table,
 - containing attributes of subclass + FK to superclass table
- object-oriented
 - each entity becomes a separate table,
 - inheriting all attributes from all superclasses
- single table with nulls
 - whole class hierarchy becomes one table,
 - containing all attributes of all subclasses (null, if unused)

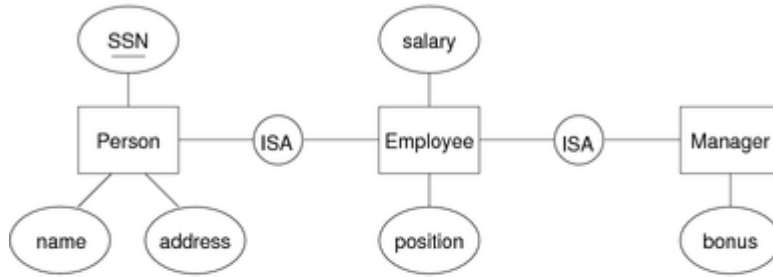
Which mapping is best depends on how data is to be used.

... Mapping Subclasses

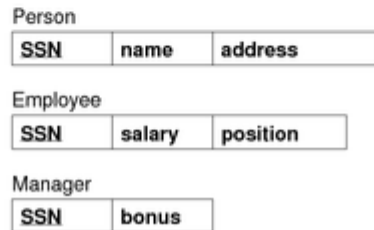
10/28

Example of ER-style mapping:

Entity-Relationship Model



Relational Model

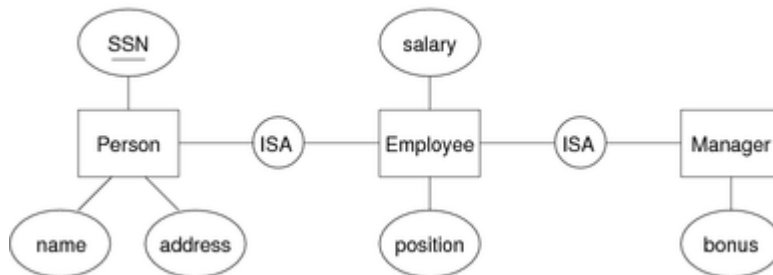


... Mapping Subclasses

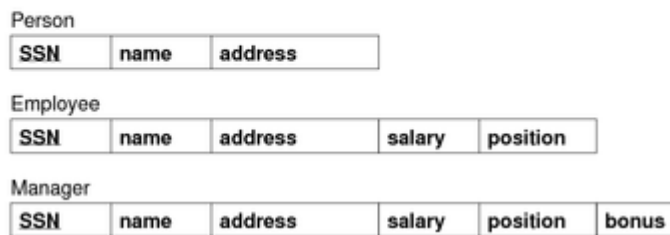
11/28

Example of object-oriented mapping:

Entity-Relationship Model



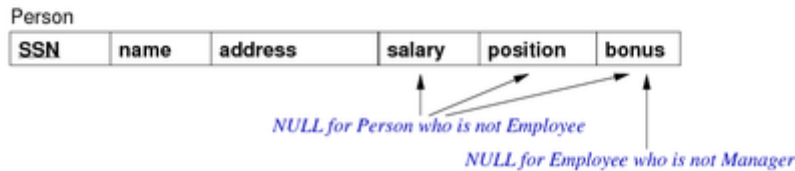
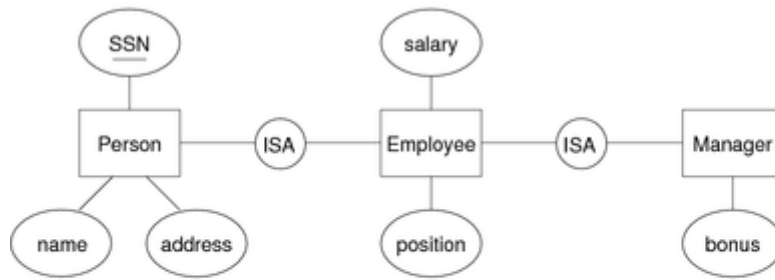
Relational Model



... Mapping Subclasses

12/28

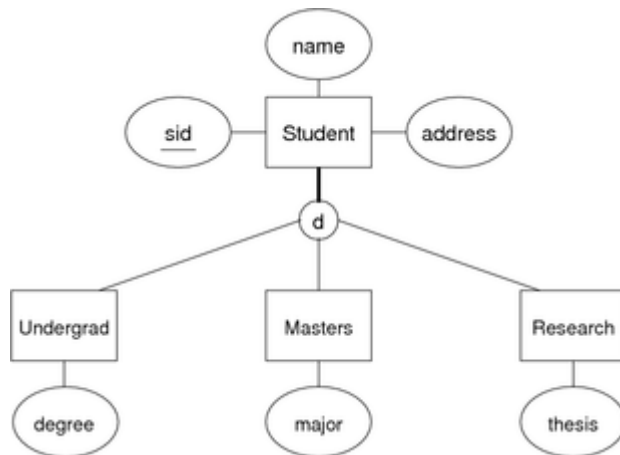
Example of single-table-with-nulls mapping:



Exercise: Disjoint subclasses

13/28

Translate the following ER design to a relational schema:



Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping

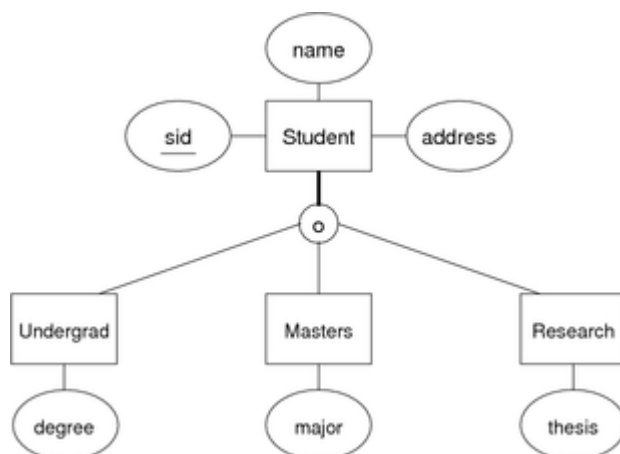
Are there aspects of the ER design that can't be mapped?

[\[Solution\]](#)

Exercise: Overlapping subclasses

14/28

Translate the following ER design to a relational schema:



Use (a) ER-mapping, (b) OO-mapping, (c) 1-table-mapping

Are there aspects of the ER design that can't be mapped?

[\[Solution\]](#)

Relational DBMSs

What is an RDBMS?

16/28

A *relational database management system* (RDBMS) is

- software designed to support large-scale data-intensive applications
- allowing high-level description of data (tables, constraints)
- with high-level access to the data (relational model, SQL)
- providing efficient storage and retrieval (disk/memory management)
- supporting multiple simultaneous users (privilege, protection)
- doing multiple simultaneous operations (transactions, concurrency)
- maintaining reliable access to the stored data (backup, recovery)

Note: databases provide *persistent* storage of information

Describing Data

17/28

RDBMSs implement \cong the relational model.

Provide facilities to define:

- domains, attributes, tuples, tables
- constraints (domain, key, referential)

Variations from the relational model:

- no strict requirement for tables to have keys
 - bag semantics, rather than set semantics
 - no standard support for general (multi-table) constraints
-

RDBMS Operations

18/28

RDBMSs typically provide at least the following:

- create/remove a database or a schema
- create/remove/alter tables within a schema
- insert/delete/update tuples within a table
- queries on data, define named queries (views)
- transactional behaviour (ACID)

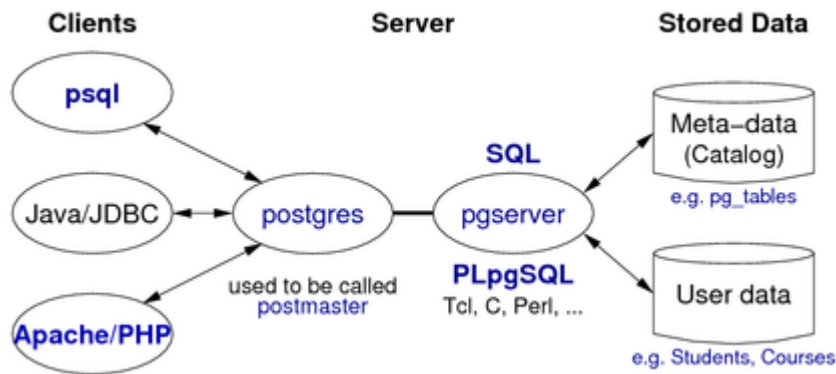
Most also provide mechanisms for

- creating/managing users of the database
 - defining/storing procedural code to manipulate data
 - implementing complex constraints (triggers)
 - defining new data types and operators (less common)
-

PostgreSQL Architecture

19/28

PostgreSQL's client-server architecture:



Using PostgreSQL

20/28

Using your PostgreSQL server in CSE (once installed):

- login to grieg, set up environment, start server
- use `psql`, etc. to manipulate databases
- stop server, log off grieg

```
wagner$ ssh YOU@grieg
grieg$ priv svr
grieg$ source /svr/YOU/env
grieg$ pg start
grieg$ psql mydb
... do stuff with your database ...
grieg$ pg stop
grieg$ exit
```

... Using PostgreSQL

21/28

PostgreSQL files (helps to understand state of server)

- PostgreSQL home directory ... `/svr/YOU/pgsql/`
- under the home directory ...
 - `postgresql.conf` ... main configuration file
 - `base/` ... subdirectories containing database files
 - `postmaster.pid` ... process ID of server process
 - `.s.PGSQL.5432` ... socket for clients to connect to server
 - `.s.PGSQL.5432.lock` ... lock file for socket
- PostgreSQL environment settings ... `/svr/YOU/env`

Building/Maintaining Databases

Managing Databases

23/28

Shell commands:

- `createdb dbname`
- `dropdb dbname`

(If no `dbname` supplied, assumes a database called `YOU`)

SQL statements:

- `CREATE DATABASE dbname`
- `DROP DATABASE dbname`

(Neither of the above is SQL-standard)

... Managing Databases

24/28

Shell commands (dump/restore):

- `pg_dump dbname > dumpfile`
- `psql dbname -f dumpfile`

(Database *dbname* is typically created just before restore)

SQL statements (used in *dumpfile*):

- `CREATE TABLE table (Attributes+Constraints)`
- `ALTER TABLE table TableSchemaChanges`
- `COPY table (AttributeNames) FROM STDIN`

Managing Tables

25/28

SQL statements:

- `ALTER TABLE table TableSchemaChanges`
- `DROP TABLE table(s) [CASCADE]`
- `TRUNCATE TABLE table(s) [CASCADE]`

(All conform to SQL standard, but all also have extensions)

`DROP . . CASCADE` drops objects which depend on the table

`TRUNCATE . . CASCADE` truncates tables which refer to the table

Managing Tuples

26/28

SQL statements:

- `INSERT INTO table (attrs) VALUES tuple(s)`
- `DELETE FROM table WHERE condition`
- `UPDATE table SET AttrValueChanges WHERE condition`

AttrValueChanges is a comma-separated list of:

- `attrname = expression`

Each list element assigns a new value to a given attribute.

Exercise: Generating IDs

27/28

Consider the following schema:

```
create table T (  
    id serial primary key,  
    x integer,  
    y varchar(10)  
);
```

- what does `serial` actually produce (look in the catalog)?
- write `INSERT` statements to add some tuples
- how could an application program get the generated id?
(`select max(id) from T` may not give the correct result; why not?)

Managing Other DB Objects

28/28

Databases contain objects other than tables and tuples:

- views, functions, sequences, types, indexes, roles, ...

Most have SQL statements for:

- `CREATE ObjectType name ...`

- **DROP** *ObjectType name ...*

Views and functions also have available:

- **CREATE OR REPLACE** *ObjectType name ...*

See PostgreSQL documentation Section IV, Chapter I for SQL statement details.
